

# Lecture 2 (UNIX System Calls part 2)

## 1 Inter Process Communication (IPC)

exit code is one way, between child and parent

## 2 Why is `fork` + `exec` good?

1. It seems wasteful to copy information of parent to just destroy the child via `exec`
2. Windows uses `create_process` instead
3. To ensure shell redirection and other features, the command requires 10 arguments
4. Also, `fork` can be made faster by just storing references and copying only when needed

## 3 More about FDT and Processes

1. Each process has its associated FDT
2. Interaction with FDT can only happen via syscalls
3. Each entry in FDT points to a certain resource and offset
4. `write` appends the offset on every successful write
5. The resource can be shared, called file sharing
6. To change the standard output, we do:

```
close(1);  
open("foo"); // smallest empty entry in FDT is at index 1
```

7. Similar as above for standard error

To write simultaneously to same file by STDOUT and STDERR, we do the following:

```
close(1);  
open("foo");  
close(2);  
dup(1); // opening foo again will not make pointers to same object, but dup will
```

The same command in shell looks like:

```
$ program > foo >2 &1
```

## 4 Pipe

1. Connecting the STDOUT of one program to STDIN of a different program
2. `pipe(int[] fdarray)` is the syntax of the syscall
3. Consider the example:

```
int fdarray[2];
pipe(fdarray);
write(fdarray[1], "hello");
read(fdarray[0], buf, 6); // "hello" is stored at buf
```

4. Notice that the pipe is created from right to left
5. pipe is used as follows

```
pipe(fdarray);
pid = fork(); // pipe is shared
if (pid > 0) {
    write(fdarray[1], ...);
} else {
    read(fdarray[0], ...);
}
```

6. Now, this is how shell implements the following pipe command: `command1 | command2`

```
int fdarray[2];
if (pipe(fdarray) < 0) panic ("error");
if ((pid = fork ()) == 0) { // child (left end of pipe)
    close (1);
    tmp = dup (fdarray[1]); // fdarray[1] is the write end, tmp will be 1
    close (fdarray[0]); // close read end
    close (fdarray[1]); // close fdarray[1]
    exec (command1, args1, 0);
} else if (pid > 0) { // parent (right end of pipe)
    close (0);
    tmp = dup (fdarray[0]); // fdarray[0] is the read end, tmp will be 0
    close (fdarray[0]);
    close (fdarray[1]); // close write end
    exec (command2, args2, 0);
} else {
    printf ("Unable to fork\n");
}
```

7. Pipe's size is limited, therefore, the commands are actually run together so that the pipe gets emptied along the way and it also helps with scheduling
8. Useful [link](#) (pipe man page)